

```

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
import statsmodels.formula.api as smf
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from google.colab import drive
drive.mount('/content/drive')
os.chdir('/MyDrive/IEOR_142_Final_Project')

```

Mounted at /content/drive

```

-----
-----
FileNotFoundError                                Traceback (most recent call
last)
<ipython-input-13-bd472d22e81a> in <module>()
      11 from google.colab import drive
      12 drive.mount('/content/drive')
--> 13 os.chdir('/MyDrive/IEOR_142_Final_Project')

FileNotFoundError: [Errno 2] No such file or directory:
'/MyDrive/IEOR_142_Final_Project'

surveydata_2014 = pd.read_csv('survey_2014.csv')
surveydata_2017 = pd.read_csv('survey_2017.csv', encoding='latin-1')
surveydata_2020 = pd.read_csv('survey_2020.csv')

surveydata_2017['Is your primary role within your company related to
tech/IT?'].value_counts()

surveydata_2020.columns.to_list()

survey2020 = surveydata_2020.drop({"#", "Is your primary role within
your company related to tech/IT?"}, axis = 1)
survey2020 = survey2020.rename({'*Are you self-employed?':
"self_employed",
    "How many employees does your company or organization have?":
"no_employees",
    "Have you ever sought treatment for a mental health disorder from a
mental health professional?": "treatment",
    "What is your gender?": "Gender",
    "What is your age?": "Age",
    "Is your employer primarily a tech company/organization?":
"tech_company",
    "Do you have a family history of mental illness?": "family_history",
    'If you have a mental health disorder, how often do you feel that it

```

```

interferes with your work *when being treated effectively?*:
"work_interface",
'Does your employer provide mental health benefits as part of
healthcare coverage?': "benefits",
"Do you know the options for mental health care available under your
employer-provided health coverage?": "care_options",
"Has your employer ever formally discussed mental health (for
example, as part of a wellness campaign or other official
communication)?": "wellness_program",
"Does your employer offer resources to learn more about mental health
disorders and options for seeking help?": "seek_help",
"Is your anonymity protected if you choose to take advantage of
mental health or substance abuse treatment resources provided by your
employer?": "anonymity",
"If a mental health issue prompted you to request a medical leave
from work, how easy or difficult would it be to ask for that leave?":
"leave",
"Would you feel comfortable discussing a mental health issue with
your direct supervisor(s)?": "supervisor",
"Would you feel comfortable discussing a mental health issue with
your coworkers?": "coworkers",
"Would you be willing to bring up a physical health issue with a
potential employer in an interview?": "phys_health_interview",
'Would you bring up your *mental* health with a potential employer in
an interview?': "mental_health_interview",
'What country do you *live* in?': "Country",
'What US state or territory do you *live* in?': "state"}, axis = 1)
survey2020 = survey2020[["self_employed", "Country", "treatment",
"Gender", "Age", "tech_company", "family_history",
"work_interface", "benefits", "care_options",
"wellness_program", "seek_help", "anonymity", "no_employees",
"leave", "supervisor", "coworkers",
"phys_health_interview", "mental_health_interview", "state"]]

```

```

survey2020['Timestamp'] = '2020'
survey2020.head(1)

```

```

survey2017 = surveydata_2017.drop({"#", "Is your primary role within
your company related to tech/IT?"}, axis = 1)
survey2017 = survey2017.rename({"<strong>Are you
self-employed?</strong>": "self_employed",
"How many employees does your company or organization have?":
"no_employees",
"Have you ever sought treatment for a mental health disorder from a
mental health professional?": "treatment",
"What is your gender?": "Gender",
"What is your age?": "Age",
"Is your employer primarily a tech company/organization?":
"tech_company",
"Do you have a family history of mental illness?": "family_history",

```

```

    "If you have a mental health disorder, how often do you feel that it
    interferes with your work <strong>when being treated
    effectively?</strong>": "work_interface",
    "Does your employer provide mental health benefitsÂ\\xa0as part of
    healthcare coverage?": "benefits",
    "Do you know the options for mental health care available under your
    employer-provided health coverage?": "care_options",
    "Has your employer ever formally discussed mental health (for
    example, as part of a wellness campaign or other official
    communication)?": "wellness_program",
    "Does your employer offer resources to learn more about mental health
    disorders and options for seeking help?": "seek_help",
    "Is your anonymity protected if you choose to take advantage of
    mental health or substance abuse treatment resources provided by your
    employer?": "anonymity",
    "If a mental health issue prompted you to request a medical leave
    from work, how easy or difficult would it be to ask for that leave?":
    "leave",
    "Would you feel comfortable discussing a mental health issue with
    your direct supervisor(s)?": "supervisor",
    "Would you feel comfortable discussing a mental health issue with
    your coworkers?": "coworkers",
    "Would you be willing to bring up a physical health issue with a
    potential employer in an interview?": "phys_health_interview",
    "Would you bring up your mental health with a potential employer in
    an interview?": "mental_health_interview",
    "What country do you <strong>live</strong> in?": "Country",
    "What US state or territory do you <strong>live</strong> in?":
    "state",
    "Submit Date (UTC)": "Timestamp"}, axis = 1)
survey2017 = survey2017[["self_employed", "Country", "treatment",
"Gender", "Age", "tech_company", "family_history",
                        "work_interface", "benefits", "care_options",
"wellness_program", "seek_help", "anonymity", "no_employees",
                        "leave", "supervisor", "coworkers",
"phys_health_interview", "mental_health_interview", "state",
"Timestamp"]]
survey2017.head(1)

survey2014 = surveydata_2014.drop({"remote_work",
"mental_health_consequence", "phys_health_consequence",
"mental_vs_physical",
                                "obs_consequence", "comments"}, axis = 1)
survey2014.head(1)

# replace some of the values that are different from the dataset from
survey2014
def replace_values(dataset):
    dataset["self_employed"] = dataset["self_employed"].replace({1:
"Yes", 0: "No"})
    dataset["treatment"] = dataset["treatment"].replace({1: "Yes", 0:

```

```

"No"})
    dataset["tech_company"] = dataset["tech_company"].replace({1: "Yes",
0: "No"})
    dataset["wellness_program"] = dataset["wellness_program"].replace("I
don't know", "Don't know")
    dataset["seek_help"] = dataset["seek_help"].replace("I don't know",
"Don't know")
    dataset["anonymity"] = dataset["anonymity"].replace("I don't know",
"Don't know")
    dataset["leave"] = dataset["leave"].replace("I don't know", "Don't
know")
    dataset["supervisor"] = dataset["supervisor"].replace("Maybe", "Some
of them")
    dataset["coworkers"] = dataset["coworkers"].replace("Maybe", "Some
of them")
    return dataset

```

```

survey2017 = replace_values(survey2017)
survey2020 = replace_values(survey2020)

```

```

# create a combined dataset

```

```

time_list = np.concatenate((survey2014[["Timestamp"]],
survey2017[["Timestamp"]], survey2020[["Timestamp"]]), axis = None)
age_list = np.concatenate((survey2014[["Age"]], survey2017[["Age"]],
survey2020[["Age"]]), axis = None)
gender_list = np.concatenate((survey2014[["Gender"]],
survey2017[["Gender"]], survey2020[["Gender"]]), axis = None)
country_list = np.concatenate((survey2014[["Country"]],
survey2017[["Country"]], survey2020[["Country"]]), axis = None)
state_list = np.concatenate((survey2014[["state"]],
survey2017[["state"]], survey2020[["state"]]), axis = None)
self_employed_list = np.concatenate((survey2014[["self_employed"]],
survey2017[["self_employed"]], survey2020[["self_employed"]]), axis =
None)
family_history_list = np.concatenate((survey2014[["family_history"]],
survey2017[["family_history"]], survey2020[["family_history"]]), axis
= None)
treatment_list = np.concatenate((survey2014[["treatment"]],
survey2017[["treatment"]], survey2020[["treatment"]]), axis = None)
work_interface_list = np.concatenate((survey2014[["work_interfere"]],
survey2017[["work_interface"]], survey2020[["work_interface"]]), axis
= None)
no_employees_list = np.concatenate((survey2014[["no_employees"]],
survey2017[["no_employees"]], survey2020[["no_employees"]]), axis =
None)
tech_company_list = np.concatenate((survey2014[["tech_company"]],
survey2017[["tech_company"]], survey2020[["tech_company"]]), axis =
None)
benefits_list = np.concatenate((survey2014[["benefits"]],
survey2017[["benefits"]], survey2020[["benefits"]]), axis = None)
care_options_list = np.concatenate((survey2014[["care_options"]],

```

```

survey2017[["care_options"]], survey2020[["care_options"]]), axis =
None)
wellness_program_list =
np.concatenate((survey2014[["wellness_program"]],
survey2017[["wellness_program"]], survey2020[["wellness_program"]]),
axis = None)
seek_help_list = np.concatenate((survey2014[["seek_help"]],
survey2017[["seek_help"]], survey2020[["seek_help"]]), axis = None)
anonymity_list = np.concatenate((survey2014[["anonymity"]],
survey2017[["anonymity"]], survey2020[["anonymity"]]), axis = None)
leave_list = np.concatenate((survey2014[["leave"]],
survey2017[["leave"]], survey2020[["leave"]]), axis = None)
supervisor_list = np.concatenate((survey2014[["supervisor"]],
survey2017[["supervisor"]], survey2020[["supervisor"]]), axis = None)
coworkers_list = np.concatenate((survey2014[["coworkers"]],
survey2017[["coworkers"]], survey2020[["coworkers"]]), axis = None)
phys_health_interview_list =
np.concatenate((survey2014[["phys_health_interview"]],
survey2017[["phys_health_interview"]],
survey2020[["phys_health_interview"]]), axis = None)
mental_health_interview_list =
np.concatenate((survey2014[["mental_health_interview"]],
survey2017[["mental_health_interview"]],
survey2020[["mental_health_interview"]]), axis = None)

```

```

df_dict = {"Timestamp": time_list, "Age": age_list, "Gender":
gender_list, "Country": country_list,
           "state": state_list, "self_employed": self_employed_list,
"family_history": family_history_list,
           "treatment": treatment_list, "work_interface":
work_interface_list, "no_employees": no_employees_list,
           "tech_company": tech_company_list, "benefits":
benefits_list, "care_options": care_options_list,
           "wellness_program": wellness_program_list, "seek_help":
seek_help_list, "anonymity": anonymity_list,
           "leave": leave_list, "supervisor": supervisor_list,
"coworkers": coworkers_list,
           "phys_health_interview": phys_health_interview_list,
           "mental_health_interview": mental_health_interview_list}

```

```

final_df = pd.DataFrame(df_dict)
final_df

```

	Timestamp	Age	...	phys_health_interview
mental_health_interview				
0	2014-08-27 11:29:31	37.0	...	Maybe
No				
1	2014-08-27 11:29:37	44.0	...	No
No				
2	2014-08-27 11:29:44	32.0	...	Yes

Yes						
3	2014-08-27	11:29:46	31.0	...		Maybe
Maybe						
4	2014-08-27	11:30:22	31.0	...		Yes
Yes						
...						...
...						
2190		2020	53.0	...		No
No						
2191		2020	23.0	...		Maybe
Maybe						
2192		2020	34.0	...		Maybe
No						
2193		2020	43.0	...		No
No						
2194		2020	37.0	...		Yes
Yes						

[2195 rows x 21 columns]

drop the unnecessary columns

final_df2 = final_df.copy()

final_df2 = final_df2.drop(["Timestamp", "state"], axis = 1)

final_df2 = final_df2.loc[~final_df2["self_employed"].isna()]

final_df3 = final_df2.dropna()

final_df3

final_df2

	Age	Gender	...	phys_health_interview	mental_health_interview
0	37.0	Female	...	Maybe	No
1	44.0	M	...	No	No
2	32.0	Male	...	Yes	Yes
3	31.0	Male	...	Maybe	Maybe
4	31.0	Male	...	Yes	Yes
...
2190	53.0	Male	...	No	No
2191	23.0	F	...	Maybe	Maybe
2192	34.0	Male	...	Maybe	No
2193	43.0	Male	...	No	No
2194	37.0	male	...	Yes	Yes

[2195 rows x 19 columns]

cleaning the gender column

gender_list = final_df2["Gender"].unique()

male_list = ["Male", "male", "M", "m", "Make", "Male ", "Cis Male",

"Male-ish", "Mal", "msle", "Mail", "cis male",

"Malr", "Cis Man", 'cis-male', 'male/androgynous ', 'cis

hetero male', 'Male (cis)',

"male (hey this is the tech industry you're talking about)", 'God King of the Valajar',

```

        'Cis male', 'man', 'Male, cis', 'cis male ', 'dude',
'maile', 'Male (CIS)', 'Man', 'Cis-male', 'cisgender male', 'MAle',
        'male', 'mail']

```

```

female_list = ["Female", "female", "Cis Female", "F", "f",
"queer/she/they", "woman", "Female ", "cis-female/femme",
        "Woman", "femal", "Female (cis)", 'Female ', "femalw",
'female (cis)', 'My sex is female.',
        'female (cisgender)', 'Female (cis) ', 'Woman-
identified', 'cis-Female', 'cis female',
        'F, cisgender', 'Female-ish', 'Trans-female', 'Femake',
'FEMALE', 'female, she/her']

```

```

other_list = ['Trans-female', 'non-binary', 'Nah', 'All', 'Enby',
'fluid', 'Genderqueer', 'Androgyne', 'Agender',
        'Guy (-ish) ^_^', 'male leaning androgynous', 'Trans
woman', 'Neuter', 'Female (trans)',
        'queer', 'ostensibly male, unsure what that really
means', "Genderfluid", "Nonbinary",
        'Non-binary', 'Agender/genderfluid', 'uhhhhhhhhhh fem
genderqueer?', 'Contextual', 'Non binary',
        'Genderqueer demigirl', 'Genderqueer/non-binary',
'nonbinary', 'trans woman', 'Transfeminine',
        'None', 'something kinda male?', 'non-binary', 'trans',
'queer', 'mostly male']

```

```

other_list2 = ["p", "A little about you", 'sometimes', '\\- ', ]
final_df2["Gender"].replace(male_list, "male", inplace = True)
final_df2["Gender"].replace(female_list, "female", inplace = True)
final_df2["Gender"].replace(other_list, "other", inplace = True)
final_df2 = final_df2.dropna(subset = ["Gender"])
final_df3 = final_df2[~final_df2["Gender"].isin(other_list2)]
final_df3

```

	Age	Gender	...	phys_health_interview	mental_health_interview
0	37.0	female	...	Maybe	No
1	44.0	male	...	No	No
2	32.0	male	...	Yes	Yes
3	31.0	male	...	Maybe	Maybe
4	31.0	male	...	Yes	Yes
...
2190	53.0	male	...	No	No
2191	23.0	female	...	Maybe	Maybe
2192	34.0	male	...	Maybe	No
2193	43.0	male	...	No	No
2194	37.0	male	...	Yes	Yes

```

[2180 rows x 19 columns]

```

```

# check Gender columns
final_df3["Gender"].unique()

array(['female', 'male', 'other'], dtype=object)

final_df4 = final_df3.copy()

labelDict = {}
for feature in final_df4:
    le = preprocessing.LabelEncoder()
    final_df4[feature] = le.fit_transform(final_df4[[feature]])

/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/
_label.py:115: DataConversionWarning: A column-vector y was passed
when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)

scaler = MinMaxScaler()
final_df4["Age"] = scaler.fit_transform(final_df4[["Age"]])
final_df4

```

	Age	Gender	...	phys_health_interview
mental_health_interview				
0	0.431034	0	...	0
1				
1	0.551724	1	...	1
1				
2	0.344828	1	...	2
2				
3	0.327586	1	...	0
0				
4	0.327586	1	...	2
2				
...
...				
2190	0.706897	1	...	1
1				
2191	0.189655	0	...	0
0				
2192	0.379310	1	...	0
1				
2193	0.534483	1	...	1
1				
2194	0.431034	1	...	2
2				

[2180 rows x 19 columns]

```

# X = final_df4[['Age', 'Gender', 'Country', 'self_employed',
'family_history', 'work_interface', 'no_employees',
# 'tech_company', 'benefits', 'care_options',

```



```

'wellness_program', 'seek_help', 'anonymity',
#         'leave', 'supervisor', 'coworkers',
'phys_health_interview', 'mental_health_interview']]
# y = final_df4["treatment"]
# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
= 0.33, random_state = 42)
train, test = train_test_split(final_df4, test_size = 0.1,
random_state = 42)

```

Logistic Regression

```

logreg = smf.logit(formula = "treatment ~ Age + Gender + Country +
self_employed + family_history + work_interface +\
no_employees + tech_company + benefits +
care_options + wellness_program + seek_help + \
anonymity + leave + supervisor + coworkers +
phys_health_interview + mental_health_interview",
data = train).fit()
print(logreg.summary)

```

Optimization terminated successfully.

Current function value: 0.591892

Iterations 5

<bound method BinaryResults.summary of

<statsmodels.discrete.discrete_model.LogitResults object at

0x7f86628af850>>

```

y_test = test['treatment']
y_prob = logreg.predict(test)
y_pred = pd.Series([1 if x > 0.5 else 0 for x in y_prob], index =
y_prob.index)

```

```

print('Accuracy: %.4f' % accuracy_score(y_test, y_pred))

```

The proportion of people that we correctly classified.

Accuracy: 0.6560

```

cm = confusion_matrix(y_test, y_pred)

```

```

print(cm)

```

```

[[58 47]
 [28 85]]

```

LDA

predicting treatment based on features

```

X_train = train.drop(columns = ['treatment'])

```

```

X_test = test.drop(columns = ['treatment'])

```

```

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

```

```

lda = LinearDiscriminantAnalysis()

```

```

y_train = train['treatment']
y_test = test['treatment']
lda = lda.fit(X_train, y_train)
lda

```

```

LinearDiscriminantAnalysis()

```

```

y_pred = lda.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
print ("Confusion Matrix: \n", cm)
#from sklearn.metrics import accuracy_score
#print("Accuracy is: \n", accuracy_score(y_test, y_pred))
acc2 = (cm.ravel()[0]+cm.ravel()[3])/sum(cm.ravel())
TPR2 = cm.ravel()[3]/(cm.ravel()[3]+cm.ravel()[2])
FPR2 = cm.ravel()[1]/(cm.ravel()[1]+cm.ravel()[0])
print('Accuracy is: %.4f' %acc2)
print('TPR is: %.4f' % TPR2)
print('FPR is: %.4f' % FPR2)

```

```

Confusion Matrix:
[[58 47]
 [28 85]]
Accuracy is: 0.6560
TPR is: 0.7522
FPR is: 0.4476

```

```

0.6559633027522935

```

```

# CART Model

```

```

from sklearn.tree import DecisionTreeClassifier

```

```

# choose optimal ccp_alpha:
from sklearn.model_selection import GridSearchCV

```

```

grid_values = {'ccp_alpha': np.linspace(0, 0.10, 201),
               'min_samples_leaf': [5],
               'min_samples_split': [20],
               'max_depth': [30],
               'class_weight' : ['balanced'],
               'random_state': [88]}

```

```

#more grid values -- will try every combo of hyper-parameters

```

```

dtc = DecisionTreeClassifier()
dtc_cv_acc = GridSearchCV(dtc, param_grid = grid_values, scoring =
'accuracy', cv=5, verbose=1)
dtc_cv_acc.fit(X_train, y_train)

```

Fitting 5 folds for each of 201 candidates, totalling 1005 fits

```
GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
             param_grid={'ccp_alpha': array([0.          , 0.0005, 0.001 ,
0.0015, 0.002 , 0.0025, 0.003 , 0.0035,
0.004 , 0.0045, 0.005 , 0.0055, 0.006 , 0.0065, 0.007 , 0.0075,
0.008 , 0.0085, 0.009 , 0.0095, 0.01  , 0.0105, 0.011 , 0.0115,
0.012 , 0.0125, 0.013 , 0.0135, 0.014 , 0.0145, 0.015 , 0.0155,
0.016 , 0.0165, 0.017 , 0.0175, 0.018 , 0.0185, 0.019 , 0.0195,
0.02  , 0.020...
0.084 , 0.0845, 0.085 , 0.0855, 0.086 , 0.0865, 0.087 , 0.0875,
0.088 , 0.0885, 0.089 , 0.0895, 0.09  , 0.0905, 0.091 , 0.0915,
0.092 , 0.0925, 0.093 , 0.0935, 0.094 , 0.0945, 0.095 , 0.0955,
0.096 , 0.0965, 0.097 , 0.0975, 0.098 , 0.0985, 0.099 , 0.0995,
0.1   ]),
             'class_weight': ['balanced'], 'max_depth':
[30],
             'min_samples_leaf': [5], 'min_samples_split':
[20],
             'random_state': [88]},
             scoring='accuracy', verbose=1)
```

```
acc = dtc_cv_acc.cv_results_['mean_test_score'] # what sklearn calls
mean_test_score is the holdout set, i.e. the validation set.
```

```
ccp = dtc_cv_acc.cv_results_['param_ccp_alpha'].data
```

```
pd.DataFrame({'ccp_alpha' : ccp, 'Validation Accuracy': acc}).head(20)
```

	ccp_alpha	Validation Accuracy
0	0	0.779318
1	0.0005	0.779318
2	0.001	0.787987
3	0.0015	0.806334
4	0.002	0.814987
5	0.0025	0.819066
6	0.003	0.816521
7	0.0035	0.816521
8	0.004	0.817031
9	0.0045	0.817031
10	0.005	0.812960
11	0.0055	0.814491
12	0.006	0.814491
13	0.0065	0.814491
14	0.007	0.814491
15	0.0075	0.814491
16	0.008	0.814491
17	0.0085	0.812960
18	0.009	0.812960
19	0.0095	0.812960

```
print('Grid best parameter ccp_alpha (max. accuracy): ',  
      dtc_cv_acc.best_params_['ccp_alpha'])  
print('Grid best score (accuracy): ', dtc_cv_acc.best_score_)
```

```
Grid best parameter ccp_alpha (max. accuracy):  0.0025  
Grid best score (accuracy):  0.819065794256634
```

```
dtc_cv_acc = dtc_cv_acc.fit(X_train, y_train)  
y_pred_cv_acc = dtc_cv_acc.best_estimator_.predict(X_test)  
cm2 = confusion_matrix(y_test, y_pred_cv_acc)  
print ("Confusion Matrix : \n", cm2)
```

```
acc2 = (cm2.ravel()[0]+cm2.ravel()[3])/sum(cm2.ravel())  
TPR2 = cm2.ravel()[3]/(cm2.ravel()[3]+cm2.ravel()[2])  
FPR2 = cm2.ravel()[1]/(cm2.ravel()[1]+cm2.ravel()[0])  
print('Accuracy is: %.4f' %acc2)  
print('TPR is: %.4f' % TPR2)  
print('FPR is: %.4f' % FPR2)
```

Fitting 5 folds for each of 201 candidates, totalling 1005 fits

Confusion Matrix :

```
[[82 23]
```

```
[21 92]]
```

Accuracy is: 0.7982

TPR is: 0.8142

FPR is: 0.2190

```
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from google.colab import drive
drive.mount('/content/drive')
os.chdir('/content/drive/MyDrive/IEOR 142 Final Project')
```

Mounted at /content/drive

```
os.listdir()
```

```
['Cipriani et al_GRISELDA_Lancet 2018_Open data.xlsx',
 'IEOR142-project-guidelines-Fall2021.pdf',
 'DepressionPredictionData',
 'Project Proposal Group 30.gdoc',
 'combined_survey.csv',
 'combined_survey2.csv',
 'Write Up Working Doc.gdoc',
 'survey14_17_20.csv',
 'survey14_17_20(2).csv',
 'Final_Project_IEOR142_Analytical_Models.ipynb',
 'tech_company_survey.csv',
 'Presentation.gslides',
 'Presentation Script.gdoc',
 'DataCleaning2.ipynb',
 'final_survey_data.csv',
 'CCP_alpha.png',
 'num_features.png',
 'bootstrap_test_accuracy.png',
 'Random_Forest_AUC.png',
 'Visualizations.ipynb',
 'Notebook.ipynb',
 'Report.gdoc',
 'Random_Forest_model.ipynb']
```

```
combined_survey = pd.read_csv('survey14_17_20.csv')
```

```
combined_survey2 = pd.read_csv('tech_company_survey.csv')
```

```
combined_survey = combined_survey.drop(['Unnamed: 0'], axis =1)
```

```
np.unique(combined_survey['treatment']).apply(lambda x: str(x))
```

```
array(['No', 'Yes', 'nan'], dtype=object)
```

```
from sklearn.model_selection import train_test_split
```

```
y = combined_survey2['treatment']
```

```
X = combined_survey2.drop(['treatment'], axis = 1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=69)
```

```
X_train.head()
```

```
      Unnamed: 0      Age  ...  phys_health_interview
mental_health_interview
118      382  0.214286  ...                        0
1
1106     1644  0.303571  ...                        0
1
486      911  0.767857  ...                        0
1
512      940  0.303571  ...                        0
1
283      611  0.535714  ...                        0
1
```

```
[5 rows x 18 columns]
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
#from sklearn.model import KFold
grid_values = {'ccp_alpha': np.linspace(0,0.2, 100),
               'min_samples_leaf': [5], # min number samples required
               'min_samples_split': [20], # min number of observations
               'max_depth': [30], # max nodes of a decision tree
               'random_state': [88]}
rfc = RandomForestClassifier()
#cv = KFold
rfc_cv_acc = GridSearchCV(rfc, param_grid = grid_values, scoring =
'accuracy', cv = 8, verbose = 1)
# note that the default metric used to optimized grid_value is
accuracy
rfc_cv_acc.fit(X_train, y_train)
```

Fitting 8 folds for each of 100 candidates, totalling 800 fits

```
GridSearchCV(cv=8, estimator=RandomForestClassifier(),
             param_grid={'ccp_alpha': array([0.00000000, 0.0020202 ,
0.0040404 , 0.00606061, 0.00808081,
0.01010101, 0.01212121, 0.01414141, 0.01616162, 0.01818182,
0.02020202, 0.02222222, 0.02424242, 0.02626263, 0.02828283,
0.03030303, 0.03232323, 0.03434343, 0.03636364, 0.03838384,
0.04040404, 0.04242424, 0.04444444, 0.04646465, 0.04848485,
0.05050505, 0.05252525, 0.05454545, 0.05656566, 0.05858586,
0.06060606, 0.06262627, 0.06464647, 0.06666667, 0.06868688,
0.07070708, 0.07272728, 0.07474749, 0.0767677 , 0.0787879 ,
0.08080811, 0.08282832, 0.08484853, 0.08686874, 0.08888895,
0.09090916, 0.09292937, 0.09494958, 0.09696979, 0.0989899 ,
0.10101011, 0.10303032, 0.10505053, 0.10707074, 0.10909095,
0.11111116, 0.11313137, 0.11515158, 0.11717179, 0.119192 ,
0.12121221, 0.12323242, 0.12525263, 0.12727284, 0.12929305,
0.13131326, 0.13333347, 0.13535368, 0.13737389, 0.1393941 ,
0.14141431, 0.14343452, 0.14545473, 0.14747494, 0.14949515,
0.15151536, 0.15353557, 0.15555578, 0.15757599, 0.1595962 ,
0.16161641, 0.16363662, 0.16565683, 0.16767704, 0.16969725,
0.17171746, 0.17373767, 0.17575788, 0.17777809, 0.1797983 ,
0.18181851, 0.18383872, 0.18585893, 0.18787914, 0.18989935,
0.19191956, 0.19393977, 0.19595998, 0.19798019, 0.19999999]),
             scoring='accuracy', cv=8, verbose=1)
```

```

0.18181818, 0.18383838, 0.18585859, 0.18787879, 0.18989899,
0.19191919, 0.19393939, 0.1959596 , 0.1979798 , 0.2      ]),
      'max_depth': [30], 'min_samples_leaf': [5],
      'min_samples_split': [20], 'random_state':
[88]],
      scoring='accuracy', verbose=1)

```

```

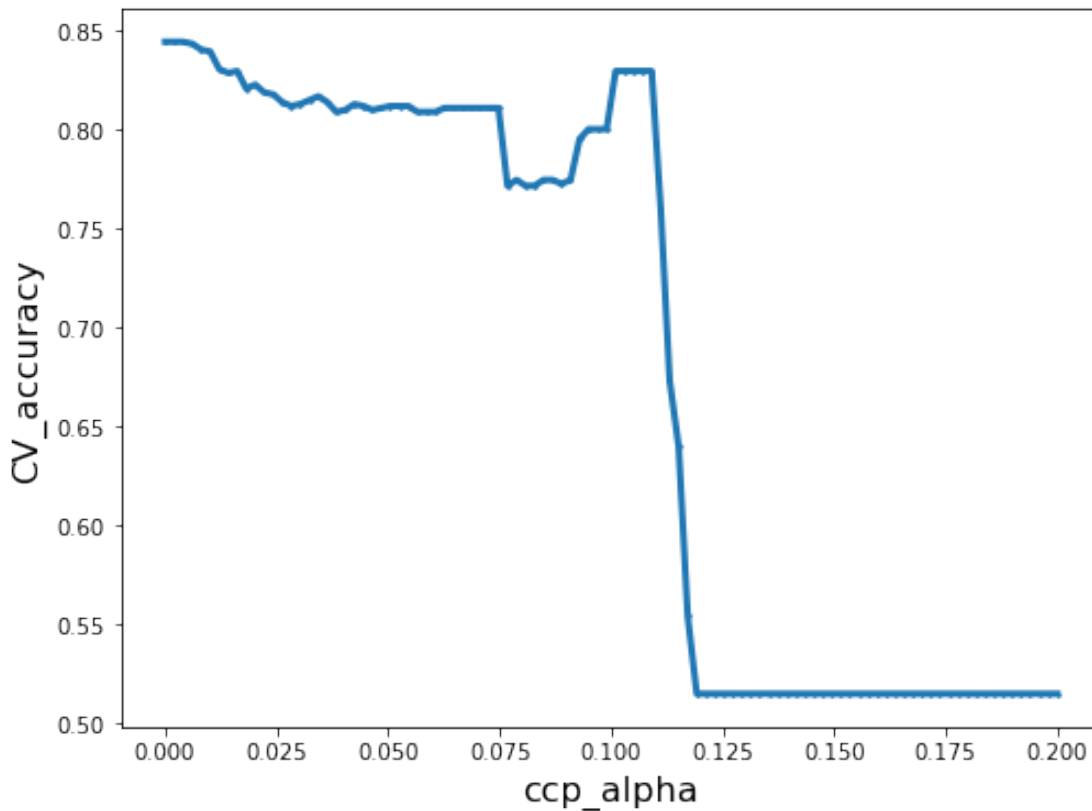
import matplotlib.pyplot as plt
ccp_alpha = rfc_cv_acc.cv_results_['param_ccp_alpha'].data
r2_score = rfc_cv_acc.cv_results_['mean_test_score']

```

```

plt.figure(figsize = (8,6))
plt.xlabel('ccp_alpha', fontsize = 16)
plt.ylabel('CV_accuracy', fontsize = 16)
plt.scatter(ccp_alpha, r2_score, s = 2)
plt.plot(ccp_alpha, r2_score, linewidth = 3)
plt.savefig('CCP alpha', dpi = 600)

```



```

print('Best CCP_alpha:', rfc_cv_acc.best_params_)
print('Grid best score Accuracy:', rfc_cv_acc.best_score_)

Best CCP_alpha: {'ccp_alpha': 0.0, 'max_depth': 30,
'min_samples_leaf': 5, 'min_samples_split': 20, 'random_state': 88}
Grid best score Accuracy: 0.8440037524606299

```

```

rfc1 = RandomForestClassifier()
rfc1.fit(X_train, y_train)
pd.DataFrame({'feature': X_train.columns, 'Importance Score':
100*rfc1.feature_importances_})

```

	feature	Importance Score
0	Unnamed: 0	9.487269
1	Age	8.218859
2	Gender	2.230060
3	Country	6.612917
4	self_employed	1.027684
5	family_history	9.505047
6	work_interface	30.253568
7	no_employees	4.204577
8	benefits	3.213617
9	care_options	4.761176
10	wellness_program	1.906289
11	seek_help	2.374829
12	anonymity	1.845887
13	leave	3.994765
14	supervisor	2.806414
15	coworkers	3.213400
16	phys_health_interview	2.505384
17	mental_health_interview	1.838258

cross validate number of columns

```

param_grid2 = {
    'max_features': list(range(1, len(X_train.columns) + 1)),
    'max_depth': [50],
    'criterion': ['entropy']}
rf2 = RandomForestClassifier()
rf_cv_acc2 = GridSearchCV(rf2, param_grid = param_grid2, scoring =
'accuracy', cv = 8, verbose = 1 )
rf_cv_acc2.fit(X_train, y_train)

```

Fitting 8 folds for each of 18 candidates, totalling 144 fits

```

GridSearchCV(cv=8, estimator=RandomForestClassifier(),
    param_grid={'criterion': ['entropy'], 'max_depth': [50],
    'max_features': [1, 2, 3, 4, 5, 6, 7, 8, 9,
10, 11, 12,
13, 14, 15, 16, 17, 18]}},
    scoring='accuracy', verbose=1)

```

```

ccp_alpha = rf_cv_acc2.cv_results_['param_max_features'].data
r2_score = rf_cv_acc2.cv_results_['mean_test_score']

```


```

plt.figure(figsize = (8,6))
plt.xlabel('Number of Features', fontsize = 16)
plt.ylabel('CV_accuracy', fontsize = 16)
plt.scatter(ccp_alpha, r2_score, s = 2)

```



```
plt.plot(ccp_alpha, r2_score, linewidth = 3)
plt.savefig('num_features', dpi = 600)
```



```
print('Grid best Max Features:',
      rf_cv_acc2.best_params_['max_features'])
```

Grid best Max Features: 7

```
final_rfc = RandomForestClassifier(max_features=
rf_cv_acc2.best_params_['max_features'],
ccp_alpha=rf_cv_acc.best_params_['ccp_alpha'])
final_rfc.fit(X_train, y_train)
final_rfc_pred = final_rfc.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
accuracy_score_1 = accuracy_score(y_test, final_rfc_pred)
print('Random Forest Classification Model', accuracy_score_1)
```

Random Forest Classification Model 0.7833001988071571

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, final_rfc_pred)
print('Confusion Matrix: \n', cm)
```

Confusion Matrix:

```
[[172  57]
 [ 52 222]]
```

y_test

```
47      1
830     0
441     0
1323    0
213     1
..
1123    0
878     0
1265    0
650     0
1109    1
```

Name: treatment, Length: 503, dtype: int64

```
tpr_rate = cm[1][1]/(cm[1][1] + cm[0][1])
fpr_rate = cm[1][0]/(cm[1][0] + cm[0][0])
```

```
print('True Positive Rate:' , tpr_rate)
print('False Positive Rate:', fpr_rate)
```

True Positive Rate: 0.7956989247311828
False Positive Rate: 0.23214285714285715

```
from sklearn.metrics import accuracy_score
def bootstrap_validation(sample_num, model, X_traindata, X_testdata,
y_traindata, y_testdata):
    model = model.fit(X_traindata, y_traindata)
    output_array = []
    for bs_iter in range(sample_num):
        bs_index = np.random.choice(X_testdata.index, len(X_testdata),
replace = True)
        bs_data = X_testdata.loc[bs_index]
        bs_label = y_testdata.loc[bs_index]
        bs_predictions = model.predict(bs_data)
        output_array.append(accuracy_score(bs_label, bs_predictions))

    return pd.DataFrame(output_array)
```

```
rfc = RandomForestClassifier(max_features=
rf_cv_acc2.best_params_['max_features'],
ccp_alpha=rfc_cv_acc.best_params_['ccp_alpha'])
bootstraps = bootstrap_validation( 50, rfc, X_train, X_test,
y_train, y_test)
```

```
CI_bootstrap = np.quantile(bootstraps, np.array([0.25, 0.75]))
CI_bootstrap
```

```

array([0.77335984, 0.79870775])

CI_list = []
for i in range(0,50):
    bootstraps = bootstrap_validation( 50, rfc, X_train, X_test,
y_train, y_test)
    CI_bootstraps = np.quantile(bootstraps,np.array([0.25,
0.75])).tolist()
    CI_list.append(CI_bootstraps)

CI_list = np.array(CI_list)
lower = CI_list[:,0]

high_bound = CI_list[:,1]

effect_measure = [sum(x)/len(x) for x in CI_list]


fig, ax = plt.subplots()
num_bins = 10
# the histogram of the data
n, bins, patches = ax.hist(effect_measure, num_bins, density=True)

mu = np.mean(effect_measure) # mean of distribution
sigma = np.std(effect_measure) # standard deviation of distribution
# add a 'best fit' line
y = ((1 / (np.sqrt(2 * np.pi) * sigma)) *
    np.exp(-0.5 * (1 / sigma * (bins - mu)**2))
ax.plot(bins, y, '--')
ax.set_xlabel('Test Accuracy')
ax.set_ylabel('Probability density')
ax.set_title('Histogram of Test Accuracy of Bootstrapped Random Forest
Models')

# Tweak spacing to prevent clipping of ylabel
plt.savefig('bootstrap_test_accuracy', dpi = 600)

```

Histogram of Test Accuracy of Bootstrapped Random Forest Models



bootstraps

```
0
0 0.769384
1 0.797217
2 0.781312
3 0.813121
4 0.797217
5 0.757455
6 0.807157
7 0.811133
8 0.763419
9 0.783300
10 0.775348
11 0.787276
12 0.815109
13 0.745527
14 0.801193
15 0.795229
16 0.787276
17 0.777336
18 0.767396
19 0.813121
20 0.823062
21 0.779324
22 0.809145
23 0.785288
24 0.801193
```

```
25 0.759443
26 0.797217
27 0.793241
28 0.795229
29 0.769384
30 0.765408
31 0.821074
32 0.789264
33 0.787276
34 0.795229
35 0.783300
36 0.795229
37 0.737575
38 0.811133
39 0.801193
40 0.799205
41 0.809145
42 0.799205
43 0.821074
44 0.763419
45 0.775348
46 0.803181
47 0.787276
48 0.785288
49 0.793241
```


```
# AUC curve
```

```
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
```

```
y_prob = final_rfc.predict_proba(X_test)[: ,1]
```

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
```

```
plt.figure(figsize=(8, 6))
plt.title('ROC Curve of Random Forest Model', fontsize=18)
plt.xlabel('FPR', fontsize=16)
plt.ylabel('TPR', fontsize=16)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, label='Random Forest (area =
{:0.2f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--',
label='Naive Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=14)
plt.show()
plt.savefig('Random_Forest_AUC', dpi = 600)
```



<Figure size 432x288 with 0 Axes>

```
plt.savefig('Random_Forest_AUC', dpi = 600)
```

<Figure size 432x288 with 0 Axes>

```

import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from google.colab import drive
drive.mount('/content/drive')
os.chdir('/content/drive/MyDrive/IEOR 142 Final Project')

```

Mounted at /content/drive

```

combined_survey2 = pd.read_csv('tech_company_survey.csv')
combined_survey = pd.read_csv('survey14_17_20.csv')

```

```
combined_survey2.head(5)
```

```

      Unnamed: 0      Age  ...  phys_health_interview
mental_health_interview
0           178  0.232143  ...                      2
1
1           179  0.303571  ...                      2
1
2           180  0.392857  ...                      0
1
3           181  0.142857  ...                      1
1
4           183  0.285714  ...                      0
1

```

[5 rows x 19 columns]

```
np.unique(combined_survey['Country'])
```

```


array(['Argentina', 'Australia', 'Austria', 'Bangladesh', 'Belarus',
      'Belgium', 'Bosnia and Herzegovina', 'Brazil', 'Bulgaria',
      'Cameroon', 'Canada', 'China', 'Colombia', 'Costa Rica',
      'Croatia',
      'Czech Republic', 'Denmark', 'Egypt', 'Finland', 'France',
      'Georgia', 'Germany', 'Greece', 'Hungary', 'Iceland', 'India',
      'Indonesia', 'Ireland', 'Israel', 'Italy', 'Japan', 'Jordan',
      'Latvia', 'Malaysia', 'Mexico', 'Moldova', 'Mongolia',
      'Netherlands', 'New Zealand', 'Nigeria', 'Norway', 'Pakistan',
      'Philippines', 'Poland', 'Portugal', 'Romania', 'Russia',
      'Serbia',
      'Singapore', 'Slovakia', 'Slovenia', 'South Africa', 'Spain',
      'Sri Lanka', 'Sweden', 'Switzerland', 'Taiwan', 'Thailand',
      'Turkey', 'Ukraine', 'United Kingdom', 'United States',
      'United States of America', 'Uruguay', 'Zimbabwe'],
      dtype=object)

```

```
# Change Country Column United States of America to United States
combined_survey['Country'] = combined_survey['Country'].apply(lambda
x: x.replace('United States of America', 'United States'))
```

```
# Create Correlation Heatplot
```

```
import seaborn as sns
plt.figure(figsize = (15,8))
corr = combined_survey2.drop(['Unnamed: 0'], axis = 1).corr()
ax = sns.heatmap(
    corr,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(20, 220, n=200),
    square=True
)
ax.set_xticklabels(
    ax.get_xticklabels(),
    rotation=45,
    horizontalalignment='right'
);
```



```
countries_count = combined_survey.groupby('Country')
['Country'].count().sort_values(ascending =
```



```
False).to_frame().rename({'Country': 'Count'}, axis = 1)
countries_count = countries_count[countries_count['Count'] >=50]


countries_list = countries_count.index
# dataset subset by countries that have a total count of 50 or more in the data
countries_data =
combined_survey[combined_survey['Country'].isin(countries_list)]

countries_data2 = countries_data.groupby(['Country',
'treatment']).agg({'treatment': 'count'}).rename({'treatment': 'Count'},
axis = 1)

countries_data2 = countries_data2.reset_index()

plt.figure(figsize = (15,8))
treatment_country = sns.barplot(x = 'Country', y = 'Count', data =
countries_data2, hue = 'treatment' , palette = "flare")
treatment_country.set_title('Seeking Mental Health Treatment by
Country')

Text(0.5, 1.0, 'Seeking Mental Health Treatment by Country')
```



```
benefits_treatment = combined_survey.groupby(['benefits',
'treatment']).agg({'treatment':
'count'}).rename({'treatment': 'Count'}, axis = 1)

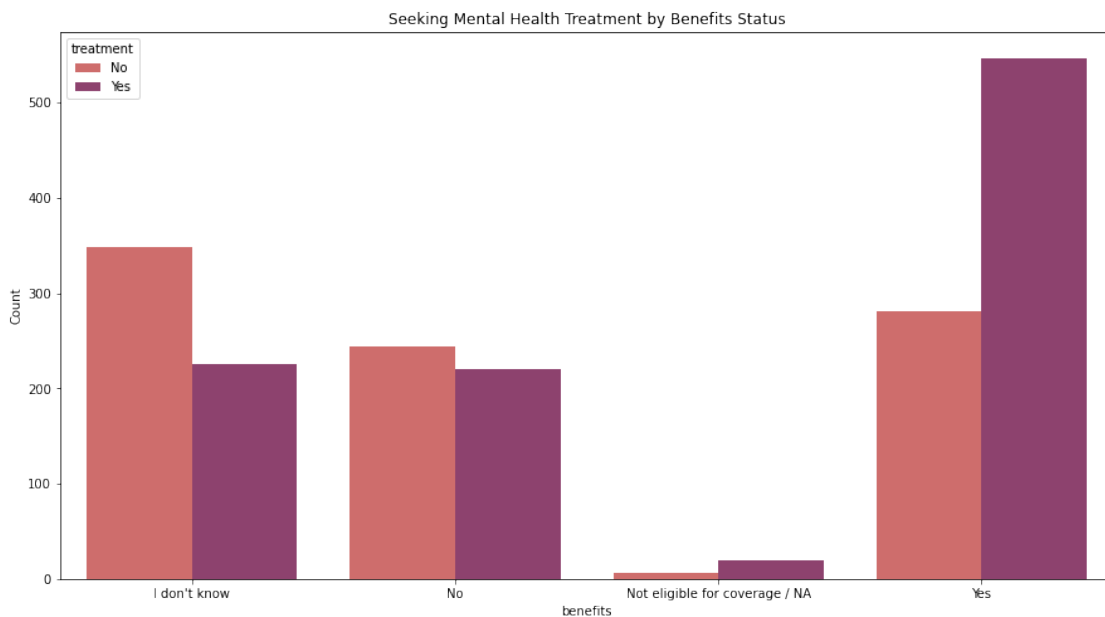
benefits_treatment = benefits_treatment.reset_index()
benefits_treatment.head()
```

	benefits	treatment	Count
0	I don't know	No	348
1	I don't know	Yes	226

2	No	No	244
3	No	Yes	220
4	Not eligible for coverage / NA	No	6

```
plt.figure(figsize = (15,8))
treatment_country = sns.barplot(x = 'benefits', y = 'Count', data =
benefits_treatment, hue = 'treatment' , palette = "flare")
treatment_country.set_title('Seeking Mental Health Treatment by
Benefits Status')
```

```
Text(0.5, 1.0, 'Seeking Mental Health Treatment by Benefits Status')
```



```
care_treatment = combined_survey.groupby(['care_options',
'treatment']).agg({'treatment':
'count'}).rename({'treatment':'Count'}, axis = 1)
care_treatment.head()
```

care_options	treatment	Count
No	No	448
	Yes	358
Not sure	No	191
	Yes	123
Yes	No	205

```
care_treatment = care_treatment.reset_index()
plt.figure(figsize = (15,8))
treatment_country = sns.barplot(x = 'care_options', y = 'Count', data =
care_treatment, hue = 'treatment' , palette = "flare")
treatment_country.set_title('Seeking Mental Health Treatment by Care
Options Status')
```

Text(0.5, 1.0, 'Seeking Mental Health Treatment by Care Options Status')

